

Article 50 Transparency Submission Pack v0.1

Short Submission Draft

Technical evidence-chain contribution to the European Commission / AI
Office consultation on draft Article 50 transparency guidelines

Kotov Ivan
Bruxelles, Belgium
2026-05-19

Status:

Draft technical submission package

Not legal advice

Not certification

Not a conformity assessment

Formal compliance remains deployment-specific and subject to legal review

Table of Contents

Article 50 Transparency Submission — Cover Note v0.1

0. Purpose
1. Submission boundary
2. Non-claim boundary
3. Context
4. Core contribution
5. Relation to the public corpus
6. Explicit bridge
7. Quiet bridge I — information theory
8. Quiet bridge II — cybernetics and control surfaces
9. Earth paragraph — relays, labels, and inspection panels
10. Expected supporting files
11. Intended use in the consultation
12. Recommended submission posture
13. References
14. Final note

Article 50 Compliance Mapping Pack v0.1

0. Executive summary
1. Non-claim boundary
2. Consultation context
3. Source corpus intake
4. Submission model
5. Role classification matrix
6. System boundary matrix
7. Article 50 obligation mapping
8. Control mapping by corpus layer
9. Provider / deployer distinction
10. CLI Agent Mesh bridge
11. Disclosure examples
12. Evidence artifacts
13. Compliance map template
14. A/B compliance language slots
15. Explicit bridge
16. Quiet bridge I — information theory
17. Quiet bridge II — Ashby / cybernetics
18. Earth paragraph — building inspection analogy
19. Recommended implementation sequence
20. Gaps and open issues
21. Recommended submission sentence
22. References
23. Final control statement

System / Role / Obligation / Control / Evidence Table v0.1

0. Executive purpose
1. Use rule
2. Evidence vocabulary
3. Responsible actor vocabulary
4. System-role table
5. Article 50 obligation-to-evidence table
6. Control-to-evidence detail table
7. System-specific table templates
8. Negative declarations table
9. EUSurvey-ready condensed table
10. Explicit bridge
11. Quiet bridge I — information theory
12. Quiet bridge II — cybernetic control
13. Earth paragraph — invoice, toolbox, and signature
14. Minimal implementation checklist
15. Red flags
16. Final table rule
17. References
18. Final control statement

Non-Claims, Open Issues, and Legal Review v0.1

0. Executive summary
1. Primary non-claim
2. Forbidden claims
3. Allowed claims
4. Legal status of this package
5. Role classification open issue
6. System boundary open issue
7. Provider / deployer open issue
8. Human review and editorial responsibility open issue
9. Machine-readable marking open issue
10. Deepfake and generated media open issue
11. Public-interest publication open issue
12. Emotion recognition open issue
13. Biometric categorisation open issue
14. Worker / enterprise deployment open issue
15. Child-facing deployment open issue
16. CLI Agent Mesh legal boundary
17. Data protection and privacy open issue
18. Security and dual-use open issue
19. Evidence sufficiency open issue
20. Submission-channel open issue
21. Publication order open issue
22. Claim strength classification
23. Review checklist
24. A/B language slots
25. Explicit bridge
26. Quiet bridge I — Ashby
27. Quiet bridge II — information theory
28. Earth paragraph — inspection boundary
29. Final non-claim statement
30. References
31. Final control rule

EUSurvey Response Draft v0.1

0. Use note
1. Respondent profile draft
2. One-sentence summary
3. Main contribution — short field
4. Main contribution — medium field
5. Main contribution — long field
6. Answer on Article 50(1): direct AI interaction disclosure
7. Answer on Article 50(2): machine-readable marking
8. Answer on Article 50(3): emotion recognition and biometric categorisation
9. Answer on Article 50(4): deepfakes and public-interest publications
10. Suggested implementation pattern for the guidelines
11. Suggested language on evidence
12. Suggested language on AI-assisted workflows and agents
13. Suggested language on public-interest AI-generated text
14. Suggested language on provenance
15. Suggested language on data minimisation
16. Suggested language on mixed roles
17. Suggested attachment note
18. Suggested closing statement
19. Ultra-short final answer
20. A/B slots for tone control
21. Claims to avoid in EUSurvey
22. Final checklist before submission
23. Recommended final response package
24. References to cite or link if allowed
25. Final control statement

Article 50 Transparency Submission — Cover Note v0.1

Author: Kotov Ivan

Location: Bruxelles, Belgium

Date: 2026-05-19

Document ID: Article50_Transparency_Cover_Note_v0_1

Package: Article50_Transparency_Submission_v0_1

Status: Draft submission cover note

Intended channel: European Commission / AI Office consultation on draft Article 50 transparency guidelines

Legal status: Not legal advice. Not certification. Not a conformity assessment.

Editorial status: Draft prepared for human review before submission.

0. Purpose

This cover note introduces a small technical submission package concerning the implementation of transparency obligations under Article 50 of the EU AI Act.

The submission does not claim that any concrete deployed system is formally compliant.

It provides an architecture-aligned contribution for mapping:

```
system
-> role
-> Article 50 obligation
-> technical control
-> evidence artifact
-> responsible actor
```

The purpose is practical: to show how transparency obligations can be expressed as an auditable engineering chain rather than as a general policy promise.

1. Submission boundary

This package is intended as a focused contribution to the Article 50 transparency guidelines consultation.

It addresses transparency controls for AI systems and AI-assisted workflows, especially where a system may:

- interact with humans;
- generate or materially assist public content;
- produce synthetic or manipulated media;
- require user-facing disclosure;
- require provenance, marking, or metadata;
- require human editorial responsibility;
- require auditability or witnessable evidence.

This package is not a request to recognize a new legal status for AI entities.

It is not a request for certification.

It is not a product declaration.

It is a technical mapping contribution.

2. Non-claim boundary

This submission does **not** claim:

- 1. that the public `c = a + b` corpus is legally compliant by itself;
- 2. that L4 Witness records automatically satisfy Article 50;
- 3. that architectural alignment equals regulatory compliance;
- 4. that any deployed system can avoid role-specific legal review;
- 5. that witness logs replace user-facing disclosure;
- 6. that metadata alone is sufficient when visible disclosure is required;
- 7. that human review is meaningful unless the responsible human actor is identified;
- 8. that open-source protocol publication is equivalent to provider or deployer compliance.

The controlling statement is:

The public corpus and related L4 Witness architecture are strongly aligned with the direction of Article 50 transparency obligations, but formal compliance depends on the concrete deployed system, role classification, user interface disclosures, marking mechanisms, audit evidence, and legal review.

3. Context

The European Commission has opened a consultation on draft guidelines for transparency obligations under Article 50 of the AI Act.

The Commission describes the guidelines as practical guidance for competent authorities, providers, and deployers of AI systems, aiming at consistent, effective, and uniform implementation of Article 50 obligations.

The Article 50 transparency rules are scheduled to apply from **2 August 2026**. They concern, among other areas, informing people when they interact with AI systems and when they are exposed to certain AI-generated or AI-manipulated content.

The Commission and AI Office are also developing supporting work on marking and labelling AI-generated content, including machine-readable marking, detection, and disclosure practices.

This submission is written for that implementation layer.

It is not a philosophical manifesto.

4. Core contribution

The submission proposes a practical compliance-mapping pattern:

Layer	Question	Required result
System boundary	What exactly is the AI system or AI-assisted workflow?	A concrete scoped system, not a general theory.
Role classification	Is the actor a provider, deployer, researcher, publisher, protocol author, or enterprise implementer?	A role-specific obligation map.

Layer	Question	Required result
Obligation trigger	Which Article 50 transparency duty is triggered?	A specific trigger, not a generic compliance claim.
Technical control	What technical or procedural control addresses the duty?	Disclosure, marking, metadata, witness, review, or routing.
Evidence artifact	What proves the control exists and was applied?	UI screenshot, log, witness record, metadata, review record, audit artifact.
Responsible actor	Who is accountable for the control and review?	Named human or organization, not an anonymous model.

This mapping is intended to reduce a common implementation failure:

policy promise
without
system boundary, role, control, evidence, or accountable actor

Article 50 transparency should not be treated as a slogan.

It should be testable.

5. Relation to the public corpus

The submission is informed by the author's public work on:

- $c = a + b$ entity architecture;
- L4 Reality Boundary;
- L4 Witness discipline;
- Actor Grounding Layer;
- Beacon / identity / attribution;
- Arbitration and Review Layer;
- CLI Agent Mesh governance;
- public release hygiene, hashes, DOI, and auditability.

These layers are not submitted as legal conclusions.

They are submitted as reusable architectural patterns for making transparency, attribution, provenance, and human responsibility more inspectable.

The central engineering claim is narrow:

A transparency obligation becomes stronger when the system can show a traceable route from user-facing notice to responsible human review and evidence.

6. Explicit bridge

$c = a + b$ places the human anchor a inside a concrete legal and operational environment.

Article 50 transparency obligations also operate through concrete actors: providers, deployers, users, publishers, editors, and affected persons.

The bridge is therefore:

```
human anchor / responsible actor
-> system role
-> transparency obligation
-> user-facing disclosure
-> witnessable evidence
```

This bridge does not turn the architecture into law.

It makes the architecture legible to compliance review.

7. Quiet bridge I — information theory

A disclosure regime is an information channel.

If the system provides too little information, the user cannot make an informed decision.

If it provides too much raw internal data, the system may expose privacy, security, trade secrets, or irrelevant cognitive noise.

Article 50 implementation therefore needs a compression discipline:

```
enough signal for the user
enough metadata for detection
enough evidence for audit
not enough raw leakage to create a new harm
```

This matches the witness principle used in the wider corpus:

```
witness the boundary, not the whole private life of the system or user
```

8. Quiet bridge II — cybernetics and control surfaces

A transparency control is a control surface.

If the notice is only decorative, it does not change the user's ability to act.

If a marking mechanism can be silently stripped, it does not preserve provenance.

If a human review checkbox has no named responsible actor, it does not create responsibility.

A useful Article 50 implementation should therefore distinguish:

- a visible user notice;
- a machine-readable mark;
- a provenance record;
- a human editorial review;
- a witness or audit record;
- a correction or dispute route.

These are different control surfaces.

They should not be collapsed into one checkbox.

9. Earth paragraph — relays, labels, and inspection panels

In a real building, a label on an electrical panel is not the same as a breaker, a breaker is not the same as an inspection report, and an inspection report is not the same as the electrician's responsibility.

All four matter.

A circuit can be dangerous even if the panel looks clean.

A panel can be technically safe but unusable if every label is wrong.

Article 50 transparency has the same structure.

A user-facing AI notice is the label.

Machine-readable provenance is the wiring map.

The audit trail is the inspection record.

The responsible human or organization is the electrician who signs the work.

A serious implementation needs all of them in the right places.

10. Expected supporting files

This cover note is intended to accompany the following package:

```
Article50_Transparency_Submission_v0_1/  
├── README.md  
├── 00_Cover_Note.md  
├── 01_Article_50_Compliance_Mapping_Pack_v0_1.md  
├── 02_System_Role_Obligation_Control_Evidence_Table_v0_1.md  
├── 03_L4_Witness_and_Provenance_Bridge_v0_1.md  
├── 04_Non_Claims_Open_Issues_and_Legal_Review_v0_1.md  
└── 05_EUSurvey_Response_Draft_v0_1.md
```

The package should remain small.

The purpose is not to submit the full public corpus.

The purpose is to submit a readable implementation bridge.

11. Intended use in the consultation

The short form of this submission is:

Article 50 transparency obligations should be mapped from concrete system boundaries and actor roles to visible disclosures, machine-readable marking, provenance metadata, human review records, and audit/witness evidence. The submitted architecture pattern shows how this can be done without confusing architectural alignment with formal legal compliance.

The proposed reading sequence is:

1. read this cover note;
2. review the compliance mapping table;
3. review the system-role-obligation-control-evidence table;
4. inspect the L4 Witness / provenance bridge;
5. check the non-claims and open issues section;

6. use the EUSurvey response draft only after human review.

12. Recommended submission posture

Use cautious wording.

Prefer:

```
architecture-aligned  
candidate evidence model  
deployment-specific  
role-dependent  
subject to legal review  
human-reviewed before submission
```

Avoid:

```
compliant  
certified  
guaranteed  
fully aligned  
proves compliance  
solves Article 50
```

The correct posture is not defensive and not promotional.

It is technical.

13. References

Official sources to be cited in the final package:

1. European Commission, consultation on the draft guidelines on transparency obligations under the AI Act, opened 8 May 2026 and closing 3 June 2026:
<https://digital-strategy.ec.europa.eu/en/consultations/consultation-draft-guidelines-transparency-obligations-under-ai-act>
 2. European Commission, draft guidelines on implementation of transparency obligations for certain AI systems under Article 50 of the AI Act:
<https://digital-strategy.ec.europa.eu/en/library/draft-guidelines-implementation-transparency-obligations-certain-ai-systems-under-article-50-ai-act>
 3. European Commission, press release on the consultation and Article 50 transparency obligations applying from 2 August 2026:
<https://digital-strategy.ec.europa.eu/en/news/commission-opens-consultation-draft-guidelines-ai-transparency-obligations>
 4. European Commission / AI Office, Code of Practice on marking and labelling of AI-generated content:
<https://digital-strategy.ec.europa.eu/en/policies/code-practice-ai-generated-content>
-

14. Final note

This submission should be understood as a compliance architecture contribution.

Its value is not that it claims final answers.

Its value is that it makes the review chain explicit:

```
system
-> role
-> obligation
-> control
-> evidence
-> responsible actor
```

That chain is the minimum useful bridge between Article 50 transparency requirements and real AI system implementation.

Article 50 Compliance Mapping Pack v0.1

Author: Kotov Ivan

Location: Bruxelles, Belgium

Date: 2026-05-19

Document ID: Article50_Compliance_Mapping_Pack_v0_1

Package: Article50_Transparency_Submission_v0_1

Status: Draft technical compliance-mapping contribution

Intended channel: European Commission / AI Office consultation on draft Article 50 transparency guidelines

Legal status: Not legal advice. Not certification. Not a conformity assessment.

Primary posture: architecture-aligned, evidence-oriented, deployment-specific, subject to legal review.

0. Executive summary

This document maps selected Article 50 transparency obligation areas to concrete technical and procedural controls from the public `c = a + b / SER / L4 / Witness / AGL / ARL / CLI Agent Mesh` corpus.

The central claim is narrow:

The corpus is strongly architecture-aligned with the direction of Article 50 transparency obligations, but formal compliance depends on the concrete deployed system, role classification, user interface disclosures, marking mechanisms, audit evidence, responsible human review, and legal review.

This document does not claim that a theory, protocol, repository, release, or AI-assisted workflow is compliant by itself.

It proposes a reviewable mapping pattern:

```
system boundary
-> actor role
-> obligation trigger
-> user-facing disclosure
-> machine-readable marking or provenance
-> witness / audit evidence
-> responsible human or organisation
```

This chain is intended to help prevent a common implementation failure:

```
general transparency language
  without
system scope, role, control, evidence, or accountable actor
```

1. Non-claim boundary

This mapping does **not** claim:

1. formal Article 50 compliance;
2. certification;
3. legal adequacy in any Member State;

4. provider or deployer status for every corpus component;
5. that open-source publication equals compliance;
6. that witness records replace user-facing disclosure;
7. that metadata alone is sufficient where visible disclosure is required;
8. that human review is meaningful without named responsibility;
9. that $c = a + b$ creates a legal person;
10. that Article 50 is the only relevant AI Act, GDPR, employment, consumer, cybersecurity, or sectoral obligation.

The intended claim is:

candidate compliance architecture;
not compliance conclusion.

2. Consultation context

The European Commission consultation on draft Article 50 transparency guidelines was published on 8 May 2026 and is open until 3 June 2026.

The stated aim is to help providers and deployers meet transparency requirements under Article 50 of the AI Act.

The Commission describes the obligations as covering, among other things:

- informing users when they interact with an AI system;
- implementing machine-readable marks in generative AI systems to enable detection of synthetic content as AI-generated or manipulated;
- informing people when they are exposed to deep fakes;
- informing people when they are exposed to AI-generated publications on matters of public interest;
- informing people about emotion recognition or biometric categorisation systems.

The transparency rules are scheduled to become applicable on 2 August 2026.

This mapping is written for that implementation layer.

3. Source corpus intake

This mapping was prepared after treating the available project material as a structured Markdown corpus, not as isolated attachments.

The available corpus archive contains 494 Markdown files across multiple repositories and corpus slices.

Relevant corpus groups include:

Corpus group	Role in this mapping
advanced-global-intelligence	Root AGI / $c = a + b$ / L4 / authorial and publication-control material.
sovereign-entity-recursion	SER, SER-FED, L4 Witness, ARL, continuity, federation, entity-governance layers.
ester-reality-bound	Actor grounding, reality-bound operational constraints, evidence discipline.

Corpus group	Role in this mapping
ester-clean-code	Executable implementation-adjacent layer and code-facing governance surface.
kot141078.github.io	Public web / publication / discoverability surface.
advanced-global-intelligence_LOCAL_ONLY	Local-only material, not treated as public submission evidence unless explicitly released.
volume-i, volume-ii, volume-iii and related book repositories	Public narrative/explanatory corpus; not primary Article 50 protocol evidence.
CCDP files provided separately in-session	Sensitive child-facing protocol layer; useful for state-not-content, handoff, conformance, but not the primary Article 50 submission surface.
c-governed-cli-agent-mesh public GitHub repository	Fresh public CLI/cloud agent governance layer, especially relevant to AI-assisted workflows, tool execution, provenance, witness, and human gates.

Important packaging note:

CCDP and child-facing materials should remain secondary in this Article 50 submission.

Reason:

Article 50 transparency
!=
full child-facing persistent AI safety.

4. Submission model

This document maps Article 50 implementation through six questions.

Step	Question	Required output
1	What is the system?	A scoped AI system, AI-assisted workflow, publication pipeline, or protocol component.
2	What is the actor role?	Provider, deployer, researcher, protocol author, enterprise implementer, publisher/ editorial actor, or mixed role.
3	What triggers the transparency obligation?	Interaction, generated content, deepfake, public-interest publication, emotion recognition, biometric categorisation, or other trigger.
4	What must the human see?	User-facing notice, publication disclosure, content label, or system explanation.
5	What must machines or auditors inspect?	Metadata, provenance, witness record, hash, review log, task contract, release manifest.
6	Who is responsible?	Named human or organisation with review, correction, escalation, or publication responsibility.

This prevents the phrase “transparent AI” from becoming a decorative label.

5. Role classification matrix

Formal obligations depend on actor role. The same author or organisation may hold different roles in different contexts.

Role	Description	Article 50 relevance	Required evidence before any compliance claim
Open-source protocol author	Publishes architectural protocols, schemas, documentation, or reference patterns.	Usually not enough by itself to create provider/deployer obligations for a concrete AI system, but may influence implementation practices.	Repository scope, license, non-claim boundary, versioned release, DOI/hash, statement that implementation-specific compliance is separate.
Researcher / independent contributor	Submits analysis, architecture patterns, or policy feedback.	Relevant as consultation input and technical contribution; not automatically a deployed-system role.	Author identity, submission scope, references, no certification claim, reproducible package.
Provider	Develops or places on the market an AI system, or has an AI system developed and places it under own name or trademark.	Potentially directly relevant to user interaction notice and machine-readable marking obligations.	Product/system definition, technical documentation, UI disclosure, marking method, evidence logs, compliance owner, legal review.
Deployer	Uses an AI system under its authority in a concrete context.	Relevant to disclosure obligations for AI interaction, deep-fakes, AI-generated public-interest publications, emotion recognition, and biometric categorisation.	Use-case scope, deployment context, user notice, publication disclosure, human review, policy, logs, legal review.
Enterprise implementer	Configures AI-assisted workflows inside an organisation, including workers, clients, documents, or operational decisions.	Relevant where users/workers/clients interact with or are affected by AI-assisted outputs.	Internal policy, worker/user notice, boundaries, task contracts, human review, data protection review, audit logs.
Publisher / editorial actor	Publishes AI-generated or AI-assisted text, images, audio, video, or reports.	Relevant for public-interest publications and generated/manipulated content disclosures.	Editorial responsibility statement, human review record, content provenance, visible label, metadata, correction procedure.
CLI/cloud agent operator	Allows AI agents to execute tasks, modify files, call tools, generate releases, or prepare publications.	Relevant as an AI-assisted workflow and provenance/witness layer; obligation depends on downstream use.	Task contract, permission scope, sandbox, witness event, human gate, review before publication, rollback record.
Child-facing system operator	Deploys persistent AI near children or educational environments.	Article 50 may apply, but additional child, education, privacy, safety, and jurisdictional duties likely dominate.	Specialist review, jurisdictional review, state-not-content disclosure, no surveillance, conformance tests, legal review.

5.1 Classification rule

Do not classify the whole corpus as one role.

Classify each concrete system or workflow separately.

Example:

```
public Markdown protocol repository
  = protocol publication / research contribution

enterprise c used in a company
  = potential deployer / enterprise implementer context

AI-assisted public article pipeline
  = publication workflow / deployer-editorial context

CLI Agent Mesh controlling executable agents
  = governance layer for AI-assisted tool execution, not automatically the public-facing AI system
```

6. System boundary matrix

System or workflow	Likely status	Article 50 relevance	Main risk of overclaim	Required next evidence
Public $c = a + b$ corpus	Protocol / research publication	Supports transparency architecture; not a deployed AI system by itself.	Claiming compliance from theory alone.	Non-claim boundary, DOI/release, references, mapping table.
L4 Witness layer	Evidence / audit / witness discipline	Supports auditability and review of privileged transitions.	Treating logs as user disclosure.	Event schema, sample witness event, privacy class, verification method.
Beacon / identity / attribution layer	Identity and actor recognition pattern	Supports distinction between entity, tool, oracle, proxy, replay, clone, and continuity-bearing system.	Treating recognition as legal status or safety proof.	Challenge process, attribution record, system role mapping.
AGL / Actor Grounding Layer	Source-state qualification	Supports grounding before reliance on source, actor, route, or claim.	Treating grounding as full legal validation.	Source-state log, liveness/provenance record, review threshold.
ARL / review / freeze / quarantine	Dispute and review layer	Supports escalation, contested claims, freeze, re-entry, and review.	Treating ARL as court or regulator.	Intake record, standing rule, decision note, human reviewer.
CLI Agent Mesh	AI-assisted executable-agent governance layer	Highly relevant to provenance of AI-assisted outputs, code, releases, reports, and submissions.	Treating agent governance as Article 50 compliance by itself.	Task contract, permission scope, witness event, human approval gate.
Ester Clean Code	Implementation-adjacent software layer	May become relevant if deployed as a system or component.	Treating repository existence as deployment evidence.	System boundary, user interaction mode, logs, UI disclosure.
Enterprise C in a company	Potential deployed AI-assisted system	Relevant if workers, clients, or third parties interact with AI or are affected by AI-generated outputs.	Under-disclosing AI role in workplace workflows.	Internal notice, policy, data flow, human review, legal/DPIA review.
AI-assisted LinkedIn / website posts	AI-assisted publication workflow	Relevant where content is AI-generated or materially AI-assisted, especially public-interest matters.	Failing to disclose material AI generation/review boundary.	Editorial note, front matter, metadata, human review record.
Generated images / audio / video / book covers	Synthetic content workflow	Relevant to generated or manipulated content marking/disclosure.	Relying only on invisible metadata.	Visible label where required, metadata, generation record, human review.
CCDP child-facing system	High-sensitivity candidate system	Article 50 may apply but is only one part of the required review.	Submitting child-safety package as generic transparency proof.	Specialist review, legal review, conformance test, jurisdictional annex.

7. Article 50 obligation mapping

This section maps major Article 50 transparency areas to candidate controls.

It does not assert final legal interpretation.

Obligation area	Trigger condition	Actor relevance	Existing architecture-aligned control	Missing deployment evidence	Required proof artifact
AI interaction disclosure	A person interacts directly with an AI system.	Provider / deployer depending on system role.	System boundary, user notice, entity/tool/oracle distinction, role labelling.	Concrete UI notice, timing, wording, user flow.	Screenshot, UI copy, interaction log, versioned policy.

Obligation area	Trigger condition	Actor relevance	Existing architecture-aligned control	Missing deployment evidence	Required proof artifact
Generative AI machine-readable marking	A generative AI system outputs AI-generated or manipulated content.	Provider primarily; deployer may rely on provider and add publication controls.	Provenance sidecar, metadata policy, hash manifest, witness reference.	Selected marking method, interoperability test, robustness limits.	Metadata sample, C2PA or equivalent record, JSON sidecar, detection test.
Deepfake disclosure	AI-generated/manipulated image, audio, or video falsely appears authentic/truthful.	Deployer / publisher / content operator.	Visible disclosure, generation record, editorial review, provenance metadata.	Exact media workflow, label placement, exception review.	Label screenshot, media metadata, review record, take-down/correction path.
AI-generated public-interest publication disclosure	AI-generated/manipulated text informs the public on matters of public interest, unless human review/editorial responsibility applies.	Deployer / publisher / editorial actor.	Human review gate, editorial responsibility declaration, publication front matter.	Whether the content is public-interest; review details; responsible human.	Editorial statement, review log, author/responsible editor, publication metadata.
Human review and editorial responsibility	Publication or workflow claims human responsibility.	Deployer / publisher.	Human gate, task contract, witness event, review outcome.	Named reviewer, review date, review scope, correction path.	Review record, signed release note, editorial responsibility statement.
Emotion recognition notice	System performs emotion recognition in covered contexts.	Provider/deployer depending use.	State classification boundaries; non-diagnostic rule; state-not-content caution.	Whether emotion recognition is actually performed; legal basis; user notice.	Negative declaration or notice, DPIA/legal review, technical description.
Biometric categorisation notice	System performs biometric categorisation in covered contexts.	Provider/deployer depending use.	Actor grounding can prevent ungrounded reliance; no default biometric claim.	Whether biometric categorisation exists; category; legal basis; notice.	Negative declaration or notice, DPIA/legal review, data flow diagram.
Provenance and attribution	Output or action needs actor/source traceability.	Provider/deployer/publisher/agent operator.	Beacon, AGL, L4 Witness, CLI Agent Mesh task records.	Concrete identifiers, event schema, retention and privacy class.	Witness event, provenance record, hash chain, source-state record.
Audit trail	Transparency control must be reviewable.	All roles when claims are made.	L4 Witness, release hygiene, SHA/DOI, task contracts, ARL notes.	Operational logging implementation.	Audit log, witness chain, release manifest, independent review where applicable.
User-accessible explanation	Users need meaningful information, not internal noise.	Provider/deployer.	Role labels, system boundary descriptions, non-claim statements.	UI text tested on actual user flow.	Notice text, UX review, version-controlled disclosure policy.

8. Control mapping by corpus layer

Corpus control	Article 50 contribution	What it can prove	What it cannot prove
$c = a + b$	Connects AI system to human anchor, responsibility, and technological substrate.	A disciplined distinction between human actor, technical system, and emergent entity relation.	Legal status, compliance, or personhood.
L4 Reality Boundary	Forces cost, time, resource, and irreversibility awareness into system design.	That the architecture avoids fantasy claims of unconstrained autonomy.	That a specific deployment is lawful.
L4 Witness	Provides tamper-evident evidence discipline for events and privileged transitions.	That a disclosure, review, generation, or action boundary can be witnessed.	That users were properly informed.
Beacon	Distinguishes entity/tool/oracle/proxy/replay/clone-style claims.	Attribution and recognition discipline.	Safety or child suitability by itself.

Corpus control	Article 50 contribution	What it can prove	What it cannot prove
Actor Grounding Layer	Grounds source, actor, route, freshness, and reliance state.	That a claim was not accepted merely because it was fluent.	Legal verification of all facts.
ARL	Routes disputes, freeze, quarantine, and review.	That contested states can stop and be reviewed.	Judicial or regulatory authority.
CLI Agent Mesh	Governs executable AI/cloud/CLI agents under task contracts and witness.	That AI-assisted operations can be bounded before file/system changes.	Article 50 compliance unless linked to a concrete user/publication workflow.
Release hygiene / SHA / DOI	Makes public artifacts stable, referenceable, and verifiable.	Integrity of released documents.	Legal adequacy of content.
Human review gate	Connects output to accountable human decision.	That a person reviewed and accepted responsibility.	That review was legally sufficient without scope and record.
Publication metadata	Provides content-level provenance and disclosure.	That a given artifact carries visible or machine-readable information.	That downstream copies preserve metadata.

9. Provider / deployer distinction

A frequent implementation error is to collapse providers and deployers.

This package recommends keeping the distinction explicit.

9.1 Provider-side transparency controls

Provider-side controls may include:

- built-in user interaction notice;
- built-in machine-readable marking for generated outputs;
- documentation for deployers;
- output provenance metadata;
- marking robustness statement;
- technical limitations;
- model/system identity disclosure;
- API output metadata;
- logs or verifiable evidence where appropriate.

9.2 Deployer-side transparency controls

Deployer-side controls may include:

- workplace/client/user notice;
- publication disclosure;
- deepfake/content label;
- human review and editorial responsibility;
- local policy;
- screenshots of UI notice;
- correction or takedown process;
- use-case-specific legal review;
- internal training and access controls.

9.3 Mixed-role cases

Many small organisations and independent implementers may act in mixed roles.

Example:

```
author builds a workflow
uses third-party LLMs
uses CLI/cloud agents
publishes AI-assisted public text
archives release artifacts
```

Possible roles:

```
deployer
publisher/editorial actor
protocol author
agent operator
```

The compliance map should not hide this.

It should show which role applies to each workflow stage.

10. CLI Agent Mesh bridge

The C-Governed CLI Agent Mesh is a useful Article 50-adjacent control layer because it addresses AI systems that do not merely generate text but may execute tasks through tools.

The relevant bridge is:

```
AI-assisted workflow
-> task contract
-> bounded permission
-> sandbox / worktree
-> action
-> witness event
-> review
-> human gate
-> publication / rejection
```

For Article 50 purposes, this helps with:

- evidence that content was AI-assisted;
- evidence that a human reviewed before publication;
- provenance of generated files;
- traceability of who authorised a release;
- correction and rollback;
- separation between agent output and human acceptance;
- no silent autonomy.

10.1 Candidate CLI Agent Mesh controls

Control	Article 50-adjacent value
Task contract	Shows what the AI agent was authorised to do.
Permission scope	Shows whether the agent could read, write, execute, publish, or access secrets.
Sandbox/worktree	Prevents uncontrolled mutation of public artifacts.
Witness event	Records the boundary event without relying on memory or chat logs.

Control	Article 50-adjacent value
Human approval gate	Separates generated output from accepted publication.
Rollback / correction path	Supports post-publication accountability.
Secrets denial	Prevents hidden private-data leakage into generated outputs.
Release manifest	Makes public artifacts inspectable and referenceable.

10.2 Non-claim

CLI Agent Mesh does not make a workflow Article 50-compliant by itself.

It provides a governance grammar for producing the evidence that a compliance review may need.

11. Disclosure examples

These are draft examples only.

They require legal and UX review before deployment.

11.1 AI interaction notice

You are interacting with an AI-assisted system. Important decisions, publications, or actions are reviewed by a responsible human operator.

11.2 AI-generated content notice

This content was generated or materially assisted by AI and reviewed by a human editor before publication.

11.3 Public-interest publication notice

This publication contains AI-assisted drafting. Human editorial responsibility remains with Kotov Ivan.

11.4 Generated media notice

This image/audio/video contains AI-generated or AI-manipulated elements.

11.5 Negative declaration for non-used sensitive functions

This system does not perform emotion recognition or biometric categorisation.

11.6 CLI-agent assisted workflow notice

This artifact was prepared through an AI-assisted workflow using bounded CLI/cloud agents. A human review gate was applied before publication.

12. Evidence artifacts

A strong Article 50 implementation should not rely on statements alone.

Evidence artifact	Purpose
UI screenshot	Shows what the user actually saw.
UI copy file	Makes notice text version-controlled.
Publication front matter	Carries disclosure inside Markdown/PDF/EPUB/web publication.
JSON sidecar metadata	Provides machine-readable provenance where embedded metadata is absent or removable.
C2PA or equivalent provenance metadata	Supports content authenticity where technically feasible.
Hash manifest	Makes released artifacts tamper-evident.
L4 Witness event	Records privileged transition or publication boundary.
CLI task contract	Shows what AI agents were authorised to do.
Human review record	Shows who reviewed, when, and with what scope.
Editorial responsibility declaration	Shows who is accountable for a public-interest publication.
Correction / takedown log	Shows feedback and remediation loop.
Legal review memo	Shows counsel or qualified review where needed.

13. Compliance map template

Every concrete implementation should be reducible to this template.

Field	Required content
System name	Name and version of the system or workflow.
System boundary	What is inside and outside the system.
Actor role	Provider / deployer / publisher / researcher / protocol author / mixed.
User class	Public user / worker / client / reader / child / internal reviewer / regulator.
Trigger	AI interaction / generated content / deepfake / public-interest publication / emotion recognition / biometric categorisation / other.
Disclosure	Exact visible notice or label.
Marking	Machine-readable metadata, watermark, C2PA, JSON sidecar, hash, or other mechanism.
Provenance	Source, model, agent, route, generation event, review event.
Human review	Reviewer, timestamp, scope, outcome.
Witness / audit	Event ID, hash, chain, privacy class, retention class.
Responsible actor	Named human or organisation.
Open legal issue	What still requires counsel or regulatory clarification.

14. A/B compliance language slots

Slot A — strict factual statement

Use when submitting to regulators or legal reviewers:

This package provides a technical mapping between Article 50 transparency obligation areas and candidate controls from the public $c = a + b$ / L4 Witness / AGL / ARL / CLI Agent Mesh corpus. It does not claim formal compliance, certification, or legal sufficiency for any concrete deployment.

Slot B — human-readable explanation

Use in public-facing explanation:

The point is not to say “this system is compliant”. The point is to show the chain: what system is used, what role the human or organisation has, what transparency duty is triggered, what notice or marking is shown, what evidence exists, and who is responsible.

If the context is legal, regulatory, or official, default to Slot A.

15. Explicit bridge

The explicit bridge is:

Article 50 transparency requires disclosure, marking, and responsibility.
 $c = a + b$ provides a disciplined separation between human anchor, technological substrate, and entity/workflow behavior.
L4 Witness, AGL, ARL, and CLI Agent Mesh provide evidence and review structures.

Therefore, the mapping chain is:

```
human anchor / responsible actor
-> AI system or workflow
-> transparency trigger
-> disclosure / marking
-> evidence
-> review / correction
```

This bridge is engineering-oriented, not metaphysical.

16. Quiet bridge I — information theory

Transparency is channel design.

A useful disclosure channel must carry enough information for the affected person to understand the AI role without forcing disclosure of irrelevant raw internals.

A useful provenance channel must carry enough information for audit and detection without creating a new privacy leak.

Therefore, Article 50 implementation should distinguish:

human-facing notice
machine-facing mark
auditor-facing evidence
operator-facing review record

These are not interchangeable.

17. Quiet bridge II — Ashby / cybernetics

A single label cannot control all failure modes.

A system that can interact, generate, publish, execute, and revise needs matching control variety:

notice for interaction
marking for synthetic output
provenance for content history
witness for privileged transitions
human review for responsibility
dispute/correction for feedback

This is an application of requisite variety.

The transparency mechanism must have enough variety to match the system's modes of action.

18. Earth paragraph — building inspection analogy

A building permit, a breaker label, an inspection report, and an electrician's signature are different things.

A breaker label helps the person in front of the panel.

The wiring diagram helps the technician.

The inspection report helps the authority.

The signature fixes responsibility.

AI transparency needs the same separation.

A user notice is the breaker label.

Machine-readable provenance is the wiring diagram.

A witness/audit event is the inspection record.

Human editorial responsibility is the signature.

If these are collapsed into one vague sentence, the system may look documented while remaining practically unauditabile.

19. Recommended implementation sequence

For any concrete system or workflow:

1. Define the system boundary.
2. Classify the actor role.
3. Identify Article 50 trigger conditions.

4. Draft visible user-facing disclosure.
5. Select machine-readable marking or metadata method where applicable.
6. Define provenance and witness events.
7. Add human review and editorial responsibility gate.
8. Define correction, takedown, dispute, or rollback path.
9. Collect screenshots, logs, metadata samples, and review records.
10. Send the package to qualified legal review.

No step should be skipped by saying “the architecture already supports transparency.”

Architecture support is not evidence.

Evidence is evidence.

20. Gaps and open issues

Current open issues for this package:

ID	Gap	Status
A50-OI-001	No final legal review yet.	Open
A50-OI-002	No final system-specific UI disclosure screenshots.	Open
A50-OI-003	No selected production marking method for all media types.	Open
A50-OI-004	No formal deployment role classification for each future system.	Open
A50-OI-005	No DPIA / GDPR review for enterprise or worker-facing deployment.	Open
A50-OI-006	No child-specific legal or developmental review for CCDP deployment.	Open
A50-OI-007	No final EUSurvey submission text accepted by human author.	Open
A50-OI-008	No final PDF / GitHub release / Zenodo archive for this Article 50 package.	Open

21. Recommended submission sentence

For EUSurvey or accompanying material:

This submission proposes a practical evidence-chain approach to Article 50 transparency implementation. It maps concrete system boundaries and actor roles to user-facing disclosures, machine-readable marking or provenance, witness/audit records, human review, and responsible actors. The approach is derived from the public $c = a + b$ / L4 Witness / AGL / ARL / C-Governed CLI Agent Mesh corpus, but it does not claim formal compliance for any deployment without system-specific evidence and legal review.

22. References

Official sources:

1. European Commission, consultation on the draft guidelines on transparency obligations under the AI Act, published 8 May 2026, open until 3 June 2026:
<https://digital-strategy.ec.europa.eu/en/consultations/consultation-draft-guidelines-transparency-obligations-under-ai-act>
2. European Commission, draft guidelines on implementation of transparency obligations for certain AI systems under Article 50 of the AI Act:
<https://digital-strategy.ec.europa.eu/en/library/draft-guidelines-implementation-transparency-obligations-certain-ai-systems-under-article-50-ai-act>
3. European Commission / AI Office, Code of Practice on marking and labelling of AI-generated content:
<https://digital-strategy.ec.europa.eu/en/policies/code-practice-ai-generated-content>

Public technical source:

1. Kotov Ivan, C-Governed CLI Agent Mesh v0.1.1, GitHub repository:
<https://github.com/Kot141078/c-governed-cli-agent-mesh>
2. Kotov Ivan, C-Governed CLI Agent Mesh v0.1.1, Zenodo DOI:
<https://doi.org/10.5281/zenodo.20257232>

23. Final control statement

The final control statement for this mapping is:

No transparency claim should stand without:

- a system boundary;
- an actor role;
- a trigger;
- a disclosure or marking control;
- an evidence artifact;
- a responsible human or organisation;
- and an open path for review, correction, or legal handoff.

This is the proposed compliance architecture bridge for Article 50.

System / Role / Obligation / Control / Evidence Table v0.1

Author: Kotov Ivan

Location: Bruxelles, Belgium

Date: 2026-05-19

Document ID: Article50_System_Role_Obligation_Control_Evidence_Table_v0_1

Package: Article50_Transparency_Submission_v0_1

Status: Draft technical evidence table

Intended channel: European Commission / AI Office consultation on draft Article 50 transparency guidelines

Legal status: Not legal advice. Not certification. Not a conformity assessment.

Primary posture: system-boundary first; role-specific; evidence-oriented; deployment-specific; subject to legal review.

0. Executive purpose

This document is the table-first operational core of the Article 50 submission package.

It translates broad transparency obligations into a repeatable engineering chain:

```
system
-> role
-> obligation trigger
-> control
-> evidence artifact
-> responsible actor
-> legal review item
```

It is designed for review by policy, legal, compliance, and technical readers.

The document intentionally avoids claiming that any existing corpus artifact is compliant by itself.

It answers a narrower question:

If a concrete system or AI-assisted workflow were reviewed under Article 50 transparency obligations, what evidence chain should be shown?

1. Use rule

Use this table only after a concrete system or workflow is named.

Do not use it to claim corpus-level compliance.

The minimum input is:

```
system name
system boundary
actor role
user or affected person class
AI interaction or generated-content trigger
publication / deployment context
```

Without those inputs, the result is:

2. Evidence vocabulary

Evidence class	Meaning	Example
EV-DISCLOSURE	Human-visible disclosure	UI banner, modal, post notice, document front matter.
EV-MARK	Machine-readable mark	C2PA record, embedded metadata, watermark, detectable marker.
EV-META	Documentary metadata	JSON sidecar, YAML front matter, provenance header.
EV-WITNESS	Tamper-evident event record	L4 Witness event, signed event, hash-linked event.
EV-CONTRACT	AI-agent task contract	CLI Agent Mesh contract, bounded scope, permissions.
EV-LOG	Operational log	Generation log, user-notice log, tool-call log.
EV-HUMAN-REVIEW	Human review evidence	Reviewer, date, scope, outcome, editorial responsibility.
EV-RELEASE	Release integrity evidence	SHA256SUMS, tag, DOI, release notes, manifest.
EV-POLICY	Policy or procedure	Disclosure policy, takedown procedure, internal AI use policy.
EV-LEGAL	Legal review artifact	Counsel memo, DPIA, role classification memo, jurisdictional review.
EV-NEGATIVE	Negative functional declaration	"No emotion recognition", "No biometric categorisation", "No autonomous publication".
EV-DRILL	Test or simulation record	Controlled scenario, red-team drill, replay, conformance test.

3. Responsible actor vocabulary

Actor	Meaning
AUTHOR	Kotov Ivan as author of the submitted package or public publication.
EDITOR	Human editor responsible for publication acceptance.
PROTOCOL_AUTHOR	Author of architecture, protocol, or schema.
PROVIDER	Actor responsible for provider-side AI system obligations, where applicable.
DEPLOYER	Actor using an AI system under its authority, where applicable.
ENTERPRISE_OPERATOR	Organisation or business operating a concrete workflow.
AGENT_OPERATOR	Human or organisation authorising CLI/cloud agents.
LEGAL_REVIEWER	Counsel or qualified legal/compliance reviewer.

Actor	Meaning
AUDITOR	Internal or external technical reviewer.
USER	Natural person interacting with or exposed to the AI system/output.

4. System-role table

System / workflow	Boundary	Likely role(s)	Article 50 relevance	Primary evidence needed	Compliance posture
Public $c = a + b$ corpus	Published Mark-down/PDF/DOI protocols and explanatory materials.	PRO- TOCOL_AUTHOR, RESEARCHER	Indirect: provides architectural transparency patterns.	EV-RELEASE, EV-META, non-claim boundary.	Architecture-aligned; not a deployed AI system by itself.
L4 Witness layer	Witness / audit / event evidence discipline.	PRO- TOCOL_AUTHOR, possible PROVIDER / DE- PLOYER component if implemented.	Supports auditability and traceability.	EV-WITNESS, schema, sample event, privacy class.	Candidate evidence control; not user disclosure by itself.
Beacon / identity / attribution layer	Recognition and attribution of entity/tool/oracle/proxy/replay-style claims.	PRO- TOCOL_AUTHOR, implementation-dependent.	Supports actor/source attribution.	Beacon record, challenge record, attribution map.	Supports transparency; not compliance alone.
Actor Grounding Layer / AGL	Source-state qualification before reliance.	PRO- TOCOL_AUTHOR, implementation-dependent.	Helps prevent ungrounded AI/source claims being relied upon.	Source-state log, route/liveness record.	Candidate control for provenance and reliance.
ARL / review / freeze / quarantine	Dispute and review routing.	PRO- TOCOL_AUTHOR, implementation-dependent.	Supports correction, contestability, and review.	Review intake, freeze record, outcome note.	Review layer; not a regulator or court.
C-Governed CLI Agent Mesh	Governance of executable CLI/cloud agents under persistent C governance.	PRO- TOCOL_AUTHOR, AGENT_OPERATOR; possibly deployer component if used.	Relevant to AI-assisted workflows and generated artifacts.	EV-CONTRACT, EV-WITNESS, sandbox record, human gate.	Strong evidence bridge; not Article 50 compliance by itself.
AI-assisted public post workflow	Drafting, editing, image generation, publication.	AUTHOR, EDITOR, possibly DEPLOYER / publisher.	May trigger AI-generated public-interest disclosure or generated media labeling.	Visible disclosure, human review record, metadata.	Deployment/publication-specific.
AI-assisted technical submission	Use of AI to draft regulatory/policy documents.	AUTHOR, EDITOR, possibly AGENT_OPERATOR.	Transparency relevance if published or submitted with AI assistance.	Editorial responsibility, provenance note, task logs.	Requires clear human responsibility.
Generated image workflow	AI-generated or manipulated images.	DEPLOYER, AUTHOR, EDITOR.	Deepfake/synthetic content marking/disclosure where applicable.	Visible label, metadata, generation record.	Depends on content type and authenticity risk.
Enterprise C for business operations	Internal worker/client/document workflow.	ENTER- PRISE_OPERATOR, possible DEPLOYER.	User/worker/client transparency may be relevant.	Internal notice, policy, logs, legal review.	Requires employment/GDPR/legal review.
Child-facing CCDP system	Persistent AI near a child, school, toy, or generated media system.	DEPLOYER, possibly PROVIDER; high-sensitivity.	Article 50 may apply but is not sufficient.	Specialist review, legal review, conformance tests, state-not-content evidence.	Not suitable as primary Article 50 submission surface.

5. Article 50 obligation-to-evidence table

ID	Obligation area	Trigger	Role	Control	Evidence artifact	Responsible actor	Legal review item
A50-001	AI interaction disclosure	Natural person interacts with an AI system.	PROVIDER / DEPLOYER	Clear notice before or during interaction.	EV-DIS-CLOSURE , screenshot, UI copy version.	PROVIDER / DEPLOYER	Whether system is in scope and notice timing is sufficient.
A50-002	Generative AI machine-readable marking	AI system generates or manipulates audio, image, video, or text.	PROVIDER	Machine-readable mark or provenance mechanism.	EV-MARK , EV-META , detection sample.	PROVIDER	Whether method is effective, interoperable, robust, technically feasible.
A50-003	Deepfake disclosure	Generated/manipulated image/audio/video resembles authentic persons, objects, places, entities, or events.	DEPLOYER / publisher	Visible disclosure to exposed persons or audience.	EV-DIS-CLOSURE , media metadata, review record.	DEPLOYER , EDITOR	Whether exceptions apply; whether content qualifies as deepfake.
A50-004	AI-generated public-interest publication disclosure	AI-generated/manipulated text informs public on matters of public interest.	DEPLOYER / publisher	Disclosure unless human review and editorial responsibility apply.	Publication note, EV-HUMAN-REVIEW , editorial statement.	EDITOR , AUTHOR	Whether publication is public-interest; sufficiency of review/responsibility.
A50-005	Human review / editorial responsibility	Human review is used as part of the disclosure logic.	EDITOR , AUTHOR , DEPLOYER	Named reviewer, scope, outcome, correction route.	EV-HUMAN-REVIEW , EV-POLICY , release note.	EDITOR , AUTHOR	Whether review process is real, documented, and sufficient.
A50-006	Emotion recognition notice	System performs emotion recognition where Article 50 applies.	PROVIDER / DEPLOYER	Explicit notice or negative declaration if absent.	EV-DIS-CLOSURE or EV-NEGATIVE , DPIA/legal review.	PROVIDER / DEPLOYER	Whether functionality counts as emotion recognition.
A50-007	Biometric categorisation notice	System performs biometric categorisation where Article 50 applies.	PROVIDER / DEPLOYER	Explicit notice or negative declaration if absent.	EV-DIS-CLOSURE or EV-NEGATIVE , data flow.	PROVIDER / DEPLOYER	Whether biometric categorisation exists and is lawful.
A50-008	User-accessible clarity	Disclosure must be understandable to natural persons.	All relevant actors	Short plain-language notice, not raw internal logs.	UI text, UX/legal review.	PROVIDER / DEPLOYER / EDITOR	Whether wording is clear, timely, and non-misleading.
A50-009	Provenance support	Generated output or AI-assisted publication requires traceability.	Mixed	Metadata, witness, source-state, task record.	EV-META , EV-WITNESS , EV-CONTRACT .	AGENT_OPERATOR , EDITOR	Retention, privacy, and disclosure boundaries.
A50-010	Correction and dispute path	Disclosure, marking, or provenance is contested or wrong.	Mixed	Correction/takedown/review workflow.	EV-POLICY , correction log, ARL-style review.	EDITOR , DEPLOYER , LEGAL_REVIEWER	Required correction timing and legal duty.

6. Control-to-evidence detail table

Control	Minimum acceptable evidence	Stronger evidence	Failure mode if missing
AI interaction notice	Screenshot of notice in user flow.	Versioned UI copy + log showing notice displayed.	User may not know they interact with AI.
Publication disclosure	Visible note in post/document/page.	Front matter + rendered page + archive hash.	Public cannot distinguish human-only from AI-assisted content.
Machine-readable mark	Embedded metadata or sidecar.	C2PA or equivalent + detection test + hash.	Synthetic content may be stripped from provenance chain.
Human editorial review	Reviewer name/role + timestamp + scope.	Signed review note + release gate + correction route.	"Human reviewed" becomes decorative.
CLI task contract	Agent task file with scope and permissions.	Contract + sandbox path + witness + review gate.	Executable agent work becomes invisible.
L4 Witness event	Event ID + timestamp + actor + action + hash.	Chain-linked witness event + privacy class + resolution note.	Privileged transition cannot be audited.
Source grounding	Source/actor/route record.	AGL record with freshness, route, liveness, and reliance state.	Ungrounded AI or proxy claim may drive publication/action.
Release integrity	SHA manifest or release tag.	SHA + signed tag + DOI + release notes.	Artifact can change without trace.
Negative sensitive-function declaration	Statement of absence.	Technical configuration + audit log + legal review.	Readers may assume emotion/biometric functions are absent without proof.
Correction path	Policy text.	Correction log + owner + response time + witness record.	Wrong label/metadata cannot be repaired visibly.

7. System-specific table templates

7.1 Public protocol repository

Field	Entry
System	Public protocol repository.
Boundary	Markdown/PDF/metadata/release artifacts only.
Role	<code>PROTOCOL_AUTHOR</code> / <code>RESEARCHER</code> .
Article 50 trigger	Usually indirect unless repository itself contains AI-generated content or interactive AI functionality.
Control	Non-claim boundary, stable release, DOI/SHA, authorship, references.
Evidence	<code>EV-RELEASE</code> , <code>EV-META</code> , <code>EV-POLICY</code> .
Responsible actor	<code>AUTHOR</code> .
Legal issue	Avoid presenting protocol publication as deployed-system compliance.

7.2 AI-assisted public post or article

Field	Entry
System	AI-assisted public publication workflow.
Boundary	Drafting, editing, image generation, human review, publication.
Role	AUTHOR , EDITOR , possible DEPLOYER .
Article 50 trigger	AI-generated/manipulated public-interest text or generated media, depending content.
Control	Visible AI assistance note where required; human editorial responsibility.
Evidence	Publication front matter, review record, metadata, generated media label.
Responsible actor	AUTHOR / EDITOR .
Legal issue	Determine whether the publication concerns matters of public interest and whether human review/editorial responsibility exception applies.

7.3 AI-generated image

Field	Entry
System	AI-generated image workflow.
Boundary	Prompt, model/tool, generated file, editing, publication.
Role	AUTHOR , DEPLOYER , possible publisher.
Article 50 trigger	Synthetic or manipulated image; deepfake if it falsely appears authentic/truthful.
Control	Visible label and/or metadata depending context.
Evidence	Generation record, metadata, visible label, review note.
Responsible actor	AUTHOR / EDITOR .
Legal issue	Whether it resembles real persons/places/events and whether disclosure must be visible.

7.4 CLI-agent assisted release

Field	Entry
System	CLI/cloud agent prepares repository changes, release notes, metadata, or package.
Boundary	Task contract, worktree, file edits, validation, human review, release.
Role	AGENT_OPERATOR , AUTHOR , possibly DEPLOYER for workflow.
Article 50 trigger	Indirect unless output is public AI-generated/assisted content or AI system interaction is user-facing.
Control	Contract, permissions, sandbox, witness, human gate, release integrity.
Evidence	EV-CONTRACT , EV-WITNESS , EV-HUMAN-REVIEW , EV-RELEASE .
Responsible actor	AGENT_OPERATOR / AUTHOR .

Field	Entry
Legal issue	Distinguish internal tool execution from public transparency obligations; preserve provenance.

7.5 Enterprise worker-facing **c**

Field	Entry
System	Enterprise/work-bound C used in company operations.
Boundary	Worker interface, document workflow, safety notices, task routing, logs.
Role	ENTERPRISE_OPERATOR , possible DEPLOYER .
Article 50 trigger	Workers or clients interact with or are affected by AI-assisted system.
Control	Worker/client notice, policy, human oversight, logs, escalation.
Evidence	UI notice, policy, access logs, review records, legal review.
Responsible actor	ENTERPRISE_OPERATOR .
Legal issue	Employment law, GDPR, DPIA, workplace surveillance, data minimisation.

8. Negative declarations table

A useful submission should also say what is **not** being claimed or performed.

Declaration	Meaning	Evidence
No formal compliance claim	The package maps controls; it does not certify compliance.	Non-claim section, cover note.
No autonomous publication	AI/agent outputs require human approval before public release.	Human gate, release record.
No emotion recognition by default	State or tone adaptation is not automatically emotion recognition.	Technical description, legal review if needed.
No biometric categorisation by default	No face/voice/body categorisation unless explicitly implemented.	Technical configuration, data-flow declaration.
No child-facing deployment in this Article 50 submission	CCDP is referenced as high-sensitivity architecture, not deployed system evidence.	Scope statement.
No vendor telemetry by default	Private memory or witness records are not ordinary vendor analytics.	Data flow, policy, implementation evidence.
No legal personhood claim	c = a + b is not submitted as a legal-personhood claim.	Cover note, non-claim boundary.

9. EUSurvey-ready condensed table

This shorter table can be converted into survey answers.

Question	Submission answer
What is the contribution?	A practical evidence-chain mapping for Article 50 transparency implementation.

Question	Submission answer
What problem does it address?	Transparency claims often lack system boundary, role, control, evidence, and responsible actor.
What is proposed?	Map every transparency obligation to visible disclosure, machine-readable marking/provenance, witness/audit evidence, and human responsibility.
What technical layers support it?	$c = a + b$, L4 Witness, Beacon, AGL, ARL, C-Governed CLI Agent Mesh.
Is compliance claimed?	No. The package is architecture-aligned and deployment-specific, not a conformity assessment.
What evidence is recommended?	UI notice, metadata, witness event, task contract, human review record, release integrity, correction path.
Why is this useful?	It makes Article 50 implementation testable rather than declarative.

10. Explicit bridge

The explicit bridge is:

Article 50 asks whether people are informed and whether generated content can be marked or disclosed. The corpus asks whether actions, identities, sources, and publications can be bounded, grounded, witnessed, and reviewed. The shared implementation problem is evidence.

Therefore:

transparency without evidence = policy decoration
evidence without disclosure = internal compliance theatre
disclosure without responsibility = label without owner

A serious Article 50 implementation needs all three.

11. Quiet bridge I — information theory

The table separates human-facing and machine-facing signals.

Reason:

a user needs understandable disclosure;
a machine needs detectable structure;
an auditor needs evidence;
a legal reviewer needs responsibility and scope.

One signal cannot serve all four audiences well.

This is a channel-design problem, not a wording problem.

12. Quiet bridge II — cybernetic control

Every control in the table should close a loop.

Loop component	Article 50 analogue
Sensor	User notice, metadata, witness record.
Comparator	Human/legal review against obligation.
Actuator	Label, correction, takedown, refusal, rollback.
Memory	Audit log, release manifest, witness chain.
Responsible controller	Human or organisation.

If a control has no feedback path, it is not a control.

It is a statement.

13. Earth paragraph — invoice, toolbox, and signature

On a real job site, a worker can have good tools and still do bad work.

A label on a box does not prove the job was done.

An invoice does not prove the wiring is safe.

A signature without inspection is theatre.

The useful chain is:

```
work order
-> tool used
-> inspected result
-> responsible signature
-> correction route if something fails
```

AI transparency needs the same discipline.

The model is the tool.

The output is the work.

The disclosure is the label.

The witness record is the inspection trail.

The human review is the signature.

14. Minimal implementation checklist

For a concrete system, require:

```
[ ] System boundary defined
[ ] Actor role classified
[ ] Article 50 trigger identified
[ ] User-facing notice drafted
[ ] Machine-readable marking selected where applicable
[ ] Provenance metadata defined
[ ] Witness/audit event defined
[ ] Human review gate defined
[ ] Responsible human or organisation named
[ ] Correction / takedown / rollback route defined
[ ] Legal review item recorded
```

If any item is missing:

do not claim compliance

15. Red flags

Red flag	Why it matters
"The system is transparent because we publish our philosophy."	Publication is not user-facing disclosure.
"The system logs everything."	Logs may create surveillance and still fail to inform users.
"The metadata is enough."	Metadata may be stripped or invisible to natural persons.
"A human was involved."	Involvement is not the same as documented review and responsibility.
"The AI did not decide, it only suggested."	AI-assisted workflows may still affect public content, users, workers, or clients.
"The agent only changed files locally."	Local changes may become public through release or publication.
"The user can ask if it is AI."	Transparency should not depend on interrogation by the user.
"The repo has a DOI."	DOI supports integrity, not legal compliance.
"The protocol is safe for children."	Child-facing claims require specialist/legal review beyond Article 50.

16. Final table rule

A system should not be described as Article 50-ready unless this row can be filled:

System	Role	Trigger	Disclosure	Marking / provenance	Evidence	Responsible actor	Legal review
...	...	...	...	...	...	...	...

If the row cannot be filled, the system may still be interesting.

It is not ready for a serious transparency claim.

17. References

Official sources:

- European Commission, consultation on draft guidelines on transparency obligations under the AI Act:
<https://digital-strategy.ec.europa.eu/en/consultations/consultation-draft-guidelines-transparency-obligations-under-ai-act>
- European Commission / AI Office, Code of Practice on marking and labelling of AI-generated content:
<https://digital-strategy.ec.europa.eu/en/policies/code-practice-ai-generated-content>

Public technical source:

- Kotov Ivan, C-Governed CLI Agent Mesh v0.1.1:
<https://github.com/Kot141078/c-governed-cli-agent-mesh>
 - DOI / Zenodo archive for C-Governed CLI Agent Mesh v0.1.1:
<https://doi.org/10.5281/zenodo.20257232>
-

18. Final control statement

The operational standard proposed by this table is:

No Article 50 transparency claim without:
system boundary
role classification
trigger identification
disclosure or marking control
evidence artifact
responsible actor
legal review path

This is the minimum evidence chain for a serious transparency implementation.

Non-Claims, Open Issues, and Legal Review v0.1

Author: Kotov Ivan

Location: Bruxelles, Belgium

Date: 2026-05-19

Document ID: Article50_Non_Claims_Open_Issues_and_Legal_Review_v0_1

Package: Article50_Transparency_Submission_v0_1

Status: Draft legal-boundary and review-control document

Intended channel: European Commission / AI Office consultation on draft Article 50 transparency guidelines

Legal status: Not legal advice. Not certification. Not a conformity assessment.

Primary posture: conservative claims, explicit limits, legal handoff before deployment-specific compliance assertions.

0. Executive summary

This document defines the non-claim boundary, open issues, and legal review requirements for the Article 50 transparency submission package.

The purpose is to prevent the package from being misread as:

- a certification claim;
- a legal compliance conclusion;
- a product conformity assessment;
- an assertion that a public research corpus is automatically Article 50 compliant;
- a statement that L4 Witness, provenance metadata, or CLI Agent Mesh governance can replace legal duties.

The controlling rule is:

```
architecture alignment
!=
formal compliance
```

A serious Article 50 claim requires:

```
concrete system boundary
-> actor role
-> obligation trigger
-> disclosure / marking / review control
-> evidence
-> responsible actor
-> legal review
```

1. Primary non-claim

This submission does not claim formal Article 50 compliance for any deployed AI system.

It claims only that the public architecture and governance corpus provides candidate patterns for:

- transparency mapping;
- provenance;
- user-facing disclosure;
- machine-readable marking support;

- witness/audit evidence;
- human review;
- responsibility assignment;
- correction and review routes.

The recommended submission sentence is:

This package provides an architecture-aligned evidence-chain approach to Article 50 transparency implementation. It does not claim formal compliance, certification, or legal sufficiency for any concrete deployment without system-specific evidence and legal review.

2. Forbidden claims

The following claims must not be made in the submission, public release, README, EUSurvey response, or cover note.

Forbidden claim	Reason
"This corpus is Article 50 compliant."	A corpus is not automatically a deployed system, provider control, or deployer workflow.
"L4 Witness proves compliance."	Witness can support evidence; it does not replace legal obligations or user-facing disclosure.
"Machine-readable metadata is enough."	Humans may still require visible disclosure. Metadata may be stripped or inaccessible.
"Human review solves public-interest disclosure."	Review must be real, scoped, recorded, and tied to editorial responsibility.
"CLI Agent Mesh makes AI-assisted workflow compliant."	Agent governance provides evidence infrastructure, not legal compliance by itself.
"Publication on GitHub/Zenodo proves regulatory adequacy."	DOI and SHA support integrity and referenceability, not legal adequacy.
"The system does not need legal review because it is architecture-safe."	Legal duties are jurisdictional and use-case-specific.
"CCDP proves child-facing compliance."	Child-facing deployment requires specialist, legal, psychological, educational, and jurisdictional review.
"The architecture replaces provider/deployer classification."	Role classification is required for concrete obligations.
"The architecture creates legal personhood for <code>C</code> ."	This package is not a legal-personhood claim.

3. Allowed claims

The following claims are acceptable if kept narrow.

Allowed claim	Condition
The corpus is architecture-aligned with Article 50 transparency direction.	Must include "formal compliance is deployment-specific."
L4 Witness can support audit evidence.	Must not be presented as user notice or legal proof.
CLI Agent Mesh can support AI-assisted workflow provenance.	Must include task contract, witness, human gate, and non-certification boundary.
<code>c = a + b</code> helps separate human anchor, technological substrate, and entity/workflow behavior.	Must not claim legal personhood or regulatory approval.

Allowed claim	Condition
Beacon / AGL / ARL can support attribution, grounding, and dispute review.	Must not claim safety or compliance without implementation evidence.
A system should map role → obligation → control → evidence.	Should be presented as a proposed implementation discipline.
Metadata and provenance are useful for Article 50 implementation.	Must include visible disclosure where required.
Human review and editorial responsibility can affect disclosure handling.	Must be documented with reviewer, scope, and responsibility.

4. Legal status of this package

This package is:

```
technical policy contribution
draft compliance mapping
evidence-chain proposal
submission support material
```

This package is not:

```
legal advice
legal opinion
conformity assessment
provider documentation
deployer declaration
DPIA
GDPR compliance report
AI Act full compliance report
product certification
safety certification
child-safety certification
cybersecurity certification
```

Any later use for a concrete deployment must be reviewed separately.

5. Role classification open issue

Article 50 obligations are role-dependent.

The package cannot classify the author, repository, or future systems under one single role.

Each concrete use case must be classified separately.

Use case	Possible role	Review status
Public protocol publication	Protocol author / researcher	Low direct Article 50 deployment relevance; still useful as policy contribution.
AI-assisted public post	Author / editor / deployer-publisher context	Needs publication-specific review.
Generated media publication	Deployer / publisher / editor	Needs media-type and deepfake-risk review.
CLI-agent assisted release	Agent operator / author / publisher	Needs provenance and human gate review.
Enterprise C in business workflow	Enterprise operator / deployer	Needs legal, employment, GDPR, data-flow review.
Child-facing CCDP implementation	Provider/deployer depending implementation	High-sensitivity; needs specialist and legal review.

Use case	Possible role	Review status
Public website / diary / publications	Author / publisher / editor	Needs AI-assisted publication disclosure policy.

Open issue:

A50-LR-001: Produce role classification memo for each concrete deployment before any compliance claim.

6. System boundary open issue

The Article 50 package must not describe “the project” as one system.

The project contains many layers:

- public theory;
- public repositories;
- DOI releases;
- narrative books;
- website publications;
- AI-assisted drafting workflows;
- CLI/cloud agent governance;
- possible enterprise systems;
- possible child-facing systems;
- local/private c systems;
- future deployment-specific products.

Each has different legal status.

Open issue:

A50-LR-002: Produce separate system boundary declarations before applying Article 50 obligations.

Minimum system boundary fields:

```
system_name
version
owner/operator
public/private/internal
user class
AI interaction mode
generated content mode
human review mode
publication mode
data categories
jurisdiction
```

7. Provider / deployer open issue

A single person or organisation may act as:

- researcher;
- author;
- protocol publisher;

- deployer;
- enterprise operator;
- editorial publisher;
- agent operator.

These roles may overlap in practice.

They should not be collapsed in the submission.

Open issue:

A50-LR-003: For each workflow, determine whether the author acts as protocol author, deployer, publisher, agent operator, enterprise operator, or another legally relevant role.

Legal review required:

yes

8. Human review and editorial responsibility open issue

Human review is not a magic exception.

A serious human review record should include:

reviewer identity
review date
artifact reviewed
scope of review
what AI assistance was used
what was changed by the human
publication decision
editorial responsibility statement
correction path

Open issue:

A50-LR-004: Define standard human review record for AI-assisted public publications and regulatory submissions.

Risk:

“human-reviewed” becomes decorative if reviewer, scope, and responsibility are not recorded.

9. Machine-readable marking open issue

The package references machine-readable marking and provenance.

It does not yet select one final marking method for all content types.

Possible marking routes include:

- C2PA or equivalent provenance metadata;
- embedded metadata;
- watermarking where appropriate;

- JSON sidecar metadata;
- release manifest;
- hash-based artifact integrity;
- visible disclosure when metadata may be stripped.

Open issue:

A50-LR-005: Select marking / metadata method per artifact type: Markdown, PDF, web page, image, audio, video, code release, dataset, and generated report.

Legal and technical review required:

yes

10. Deepfake and generated media open issue

Generated images, covers, illustrations, audio, or video may require different treatment depending on whether they:

- are clearly illustrative;
- resemble real people;
- imitate a real voice;
- depict real events;
- may mislead a viewer into believing they are authentic;
- are used in public-interest context.

Open issue:

A50-LR-006: Define generated media disclosure policy and deepfake-risk review checklist.

Minimum review fields:

```
media_type
real_person_likeness
real_place_or_event_likeness
voice_or_face_imitation
public_interest_context
visible_label_required
metadata_required
human_review_completed
```

11. Public-interest publication open issue

Article 50 includes obligations around AI-generated or manipulated text published to inform the public on matters of public interest, subject to conditions and exceptions.

The package cannot determine automatically whether every article, diary entry, post, or technical document is a public-interest publication.

Open issue:

A50-LR-007: Define public-interest publication classification checklist.

Candidate criteria:

- topic concerns public policy, AI safety, regulation, rights, risk, public institutions, or social impact;
- publication targets the general public or professional public;
- text is materially AI-generated or AI-assisted;
- human editorial responsibility exists and is documented;
- disclosure is visible where required or prudent.

Legal review required:

yes

12. Emotion recognition open issue

The corpus contains emotional, state, dependency, soft safety, and affective terminology in some contexts.

That does not automatically mean the system performs legally relevant emotion recognition.

However, the boundary must be reviewed carefully.

Open issue:

A50-LR-008: Distinguish emotional style adaptation, state signaling, dependency audit, and legally relevant emotion recognition.

Negative declaration is allowed only if technically true:

This system does not perform emotion recognition.

But where the system infers or categorises emotions from natural persons, legal review is required.

Risk:

Architectural language about emotion may be misread as an implemented emotion-recognition system.

13. Biometric categorisation open issue

The Article 50 package does not claim that biometric categorisation is used.

If a future system uses:

- face analysis;
- voice identity;
- gait;
- body features;
- biometric templates;
- demographic or sensitive categorisation from biometric data;

then separate legal review is mandatory.

Open issue:

14. Worker / enterprise deployment open issue

An enterprise **c** in a construction or operational company may involve:

- workers;
- clients;
- subcontractors;
- safety instructions;
- work quality records;
- photos from sites;
- documents;
- scheduling;
- material advice;
- possible AI interaction.

This may trigger more than Article 50:

- GDPR;
- employment law;
- workplace surveillance rules;
- safety obligations;
- client confidentiality;
- liability and professional responsibility;
- sectoral rules.

Open issue:

A50-LR-010: Enterprise c deployment requires separate Belgian/EU legal review, worker notice, data-flow mapping, and internal policy.

Article 50 transparency mapping is necessary but not sufficient.

15. Child-facing deployment open issue

CCDP is not the main submission surface for this Article 50 package.

Reason:

child-facing persistent AI safety
>
Article 50 transparency

Child-facing systems require:

- child-development review;
- education review;
- legal/compliance review;
- privacy/security review;

- jurisdictional review;
- conformance tests;
- red-team review using synthetic fixtures only;
- safe release split between public, technical, restricted, and internal files.

Open issue:

A50-LR-011: Do not present CCDP as Article 50 compliance. Present it only as a high-sensitivity future extension and evidence of state-not-content, handoff, and non-surveillance discipline.

16. CLI Agent Mesh legal boundary

C-Governed CLI Agent Mesh can support provenance for AI-assisted executable workflows.

It can provide:

- task contracts;
- permission scopes;
- sandbox records;
- witness events;
- human gates;
- rollback and review;
- no silent autonomy.

It cannot by itself provide:

- Article 50 certification;
- legal compliance;
- safe deployment;
- user-facing disclosure;
- provider/deployer classification;
- sector-specific legal compliance.

Open issue:

A50-LR-012: Define how CLI Agent Mesh witness records connect to publication disclosure and human editorial responsibility records.

17. Data protection and privacy open issue

Article 50 transparency is not the whole data protection problem.

Any concrete system may also require review under:

- GDPR;
- ePrivacy where relevant;
- employment privacy rules;
- child data protection rules;
- data minimisation requirements;
- retention and deletion policy;
- cross-border transfer rules;
- vendor/cloud data processing terms.

Open issue:

A50-LR-013: Create separate data-flow and data-protection review for any deployed system.

This package does not perform that review.

18. Security and dual-use open issue

Some corpus layers involve:

- CLI/cloud agents;
- tool execution;
- sandboxing;
- secrets;
- privilege escalation prevention;
- witness/audit records;
- defensive containment;
- generated code;
- release automation.

These are legitimate governance topics.

They must not be converted into:

- offensive automation;
- malware;
- credential extraction;
- unauthorized system access;
- hack-back;
- evasion instructions;
- uncontrolled autonomous action.

Open issue:

A50-LR-014: Security review required before any executable-agent deployment or publication of operational details.

Submission posture:

defensive governance only
no offensive operationalization
no silent autonomy
no secrets exposure

19. Evidence sufficiency open issue

The package currently proposes evidence types.

It does not yet contain final evidence for a concrete system.

Open issue:

A50-LR-015: For each concrete system, collect evidence pack before claiming readiness.

Minimum evidence pack:

```
role classification memo
system boundary declaration
UI disclosure screenshot
sample generated output with disclosure
sample metadata / mark
sample witness event
human review record
editorial responsibility statement
release manifest / SHA
correction path
legal review memo
```

20. Submission-channel open issue

The official consultation channel is an online questionnaire.

This package is not submitted merely by preparing Markdown files.

Open issue:

A50-LR-016: Convert this package into a concise EUSurvey response and submit through the official consultation questionnaire before the deadline.

Supporting documents may be linked or attached only if the questionnaire permits.

The final answer must be short enough to be read.

21. Publication order open issue

Recommended publication order:

```
draft internal package
-> human review
-> legal/claim review
-> PDF generation
-> GitHub release
-> Zenodo DOI if appropriate
-> EUSurvey submission
-> archive submitted text
```

Do not publish a DOI package if the claims are not stable.

Open issue:

A50-LR-017: Decide whether this package should receive DOI before or after EUSurvey submission.

Recommended default:

GitHub first, DOI after final freeze.

22. Claim strength classification

Claim class	Meaning	Allowed now?
A50-C0	Idea / draft note.	Yes
A50-C1	Architecture-aligned mapping.	Yes
A50-C2	Candidate control proposal.	Yes
A50-C3	Evidence model proposal.	Yes
A50-C4	Implementation-ready for a named system.	Not yet
A50-C5	Legally reviewed deployment mapping.	Not yet
A50-C6	Formal compliance claim.	Not allowed in this package
A50-C7	Certification / conformity assessment.	Not allowed

Current package maximum claim:

A50-C3 – evidence model proposal

Not higher.

23. Review checklist

Before submission, check:

- ☐ No claim of formal compliance
- ☐ No claim of certification
- ☐ No claim that witness replaces disclosure
- ☐ No claim that metadata replaces visible notice
- ☐ No claim that CLI Agent Mesh itself is Article 50 compliant
- ☐ Provider/deployer distinction preserved
- ☐ Human review defined as recorded responsibility, not decoration
- ☐ CCDP kept secondary and high-sensitivity
- ☐ Legal review items clearly marked
- ☐ Official EUSurvey submission path identified
- ☐ Package is short enough for policy review
- ☐ Public links are stable
- ☐ No private or sensitive implementation details exposed

24. A/B language slots

Slot A – strict official wording

This submission does not claim formal Article 50 compliance. It provides a technical evidence-chain mapping for implementation discussion: system boundary, actor role, transparency trigger, disclosure or marking control, evidence artifact, responsible actor, and legal review path.

Slot B — plain explanation

The point is not to say “we are compliant.” The point is to show what a serious transparency claim should contain: what system is involved, who is responsible, what duty is triggered, what notice or mark is used, what evidence exists, and what still needs legal review.

Use Slot A for the consultation.

25. Explicit bridge

The explicit bridge is:

Article 50 transparency is legal and procedural.
The corpus contributes architecture and evidence discipline.
The bridge is valid only when the architecture is attached to a concrete system, role, obligation, control, evidence artifact, and responsible actor.

This prevents both underclaim and overclaim.

26. Quiet bridge I — Ashby

A system with many modes of action requires many modes of control.

A single statement cannot cover:

- AI interaction;
- text generation;
- generated images;
- deepfake risk;
- public-interest publication;
- executable agents;
- human review;
- release automation;
- enterprise workflows;
- child-facing systems.

The compliance apparatus must have enough variety to match the system variety.

This is why the package uses matrices instead of one declaration.

27. Quiet bridge II — information theory

Compliance evidence should transmit enough information to support review.

It should not transmit everything.

The correct compression is:

enough signal for legal review
enough signal for public disclosure
enough signal for technical audit
minimal raw private leakage

A good evidence pack is not a data dump.

It is a controlled channel.

28. Earth paragraph — inspection boundary

In construction, a technical drawing is not a permit.

A permit is not a finished wall.

A finished wall is not an inspection certificate.

An inspection certificate is not a guarantee that nobody will ever misuse the building.

Each document has a boundary.

The same is true here.

The architecture explains the design.

The evidence pack shows what happened.

The legal review checks obligations.

The submission gives policy input.

Mixing these layers creates false confidence.

False confidence is worse than an open issue.

29. Final non-claim statement

The final non-claim statement for the whole package is:

This Article 50 package is a technical, architecture-aligned, evidence-oriented contribution to transparency implementation. It is not legal advice, not certification, not a conformity assessment, and not a claim that any deployed system is compliant. Formal compliance requires concrete system definition, role classification, disclosure and marking implementation, evidence collection, responsible human review, and qualified legal review.

This statement should appear in the final package and in any PDF or release notes.

30. References

Official sources:

1. European Commission, consultation on draft guidelines on transparency obligations under the AI Act:
<https://digital-strategy.ec.europa.eu/en/consultations/consultation-draft-guidelines-transparency-obligations-under-ai-act>
2. European Commission, draft guidelines on implementation of transparency obligations for certain AI systems under Article 50 of the AI Act:
<https://digital-strategy.ec.europa.eu/en/library/draft-guidelines-implementation-transparency-obligations-certain-ai-systems-under-article-50-ai-act>

3. European Commission / AI Office, Code of Practice on marking and labelling of AI-generated content:
<https://digital-strategy.ec.europa.eu/en/policies/code-practice-ai-generated-content>

Public technical source:

1. Kotov Ivan, C-Governed CLI Agent Mesh v0.1.1:
<https://github.com/Kot141078/c-governed-cli-agent-mesh>
2. DOI / Zenodo archive for C-Governed CLI Agent Mesh v0.1.1:
<https://doi.org/10.5281/zenodo.20257232>

31. Final control rule

Do not move from:

architecture-aligned

to:

compliant

without:

named system
role classification
implemented disclosure
implemented marking/provenance
evidence artifact
responsible actor
legal review

That is the safety rail for the Article 50 package.

EUSurvey Response Draft v0.1

Author: Kotov Ivan

Location: Bruxelles, Belgium

Date: 2026-05-19

Document ID: Article50_EUSurvey_Response_Draft_v0_1

Package: Article50_Transparency_Submission_v0_1

Status: Draft questionnaire response support document

Intended channel: European Commission / AI Office consultation on draft Article 50 transparency guidelines

Legal status: Not legal advice. Not certification. Not a conformity assessment.

Submission status: Not submitted. Requires human review before use.

0. Use note

This document is a drafting aid for the official EUSurvey consultation response.

It is not the submission itself.

The official response must be entered through the European Commission's online consultation questionnaire before the deadline.

This draft should be adapted to:

- the actual questionnaire fields;
- word or character limits;
- available attachment/link options;
- final human review;
- legal/claim review before submission.

Do not paste the entire package into the questionnaire.

Use concise answers and link or attach the supporting package only where the questionnaire allows.

1. Respondent profile draft

Use or adapt:

```
Name: Kotov Ivan
Location: Bruxelles, Belgium
Stakeholder type: Independent researcher / independent consultant / AI systems architect / protocol author
Relevant work: c = a + b protocol, L4 Reality Boundary, L4 Witness, Actor Grounding Layer, Arbitration / Review Layer, C-Governed CLI Agent Mesh, AI transparency and evidence-chain architecture
Submission type: Technical implementation contribution / evidence-chain mapping for Article 50 transparency obligations
```

If the form asks for organisation:

```
Independent individual / independent researcher
```

If the form asks whether the response may be published:

```
To be selected by the author after final review.
```

Recommended public-facing descriptor:

Independent researcher and AI systems architect, Bruxelles, Belgium.

Avoid:

Certified compliance expert
Official legal representative
Provider of a certified Article 50 system

unless legally verified.

2. One-sentence summary

Short version:

This submission proposes an evidence-chain approach to Article 50 transparency: each obligation should be mapped from a concrete system boundary and actor role to user-facing disclosure, machine-readable marking or provenance, witness/audit evidence, human review, and a responsible actor.

Very short version:

Article 50 transparency should be implemented as a traceable chain: system, role, trigger, disclosure, marking/provenance, evidence, human responsibility.

3. Main contribution — short field

Use if the questionnaire has a short “main feedback” field.

I support practical guidance that makes Article 50 obligations operational for providers and deployers. My contribution is an evidence-chain mapping pattern: every transparency claim should identify the concrete system boundary, actor role, obligation trigger, user-facing disclosure, machine-readable marking or provenance, witness/audit evidence, human review, and responsible actor. This avoids both overclaiming and purely decorative transparency notices.

4. Main contribution — medium field

Use if the questionnaire allows a medium-length answer.

My submission proposes a practical evidence-chain approach to Article 50 transparency implementation. In my view, transparency should not be treated as a general declaration that “AI was used”. A useful transparency control should be mapped to a concrete system and role: provider, deployer, publisher/editorial actor, enterprise operator, or AI-assisted workflow operator.

For each obligation, the implementation should show: the trigger, the visible disclosure shown to the person, the machine-readable mark or provenance mechanism where applicable, the audit or witness evidence that the control was applied, the human review record where relevant, and the responsible human or organisation.

This is especially important for AI-assisted publication workflows and executable agent workflows. A model may generate text, but a tool-using AI agent may also create files, modify repositories, generate metadata, prepare release artifacts, or assist with public submissions. The guidelines should encourage a separation between AI output, provenance evidence, and human acceptance before publication.

The submission does not claim formal compliance or certification. It provides an architecture-aligned implementation pattern based on public work around $c = a + b$, L4 Witness, Actor Grounding Layer, Arbitration / Review Layer, and C-Governed CLI Agent Mesh.

5. Main contribution — long field

Use if the questionnaire allows a longer free-text response.

I welcome the Commission's effort to provide practical guidance on Article 50 transparency obligations. My main suggestion is that the guidelines should encourage providers and deployers to implement transparency as an evidence chain, not as a single policy statement or label.

A serious transparency claim should be mapped through the following chain:

system boundary -> actor role -> obligation trigger -> user-facing disclosure -> machine-readable marking or provenance -> witness/audit evidence -> human review -> responsible actor -> correction or review route.

This mapping is important because the same AI-assisted workflow can contain several distinct transparency surfaces. A user-facing notice informs a person. A machine-readable mark supports detection. A provenance record supports traceability. A witness or audit record supports review. Human editorial responsibility assigns accountability. A correction route closes the feedback loop. These functions should not be collapsed into one generic checkbox.

I recommend that the guidelines explicitly encourage implementers to preserve the provider/deployer distinction while also addressing mixed-role situations. Independent creators, SMEs, researchers, publishers, and enterprise operators may use third-party models, local tools, and CLI/cloud agents in the same workflow. In such cases, the relevant question is not only “which model generated the output?” but also “who deployed the system, who reviewed the output, who accepted responsibility, and what evidence exists?”

This is especially relevant for AI-assisted public-interest publications and generated media. Where a human review or editorial responsibility path is used, it should be substantive and evidenced. A mere statement that “a human was involved” should not be enough. A useful record should identify the reviewer or responsible editor, the reviewed artifact, the review scope, the decision, the timestamp, and the correction route.

The same reasoning applies to executable agent workflows. AI agents may not only generate text; they may create files, modify code, run scripts, prepare metadata, update releases, or publish artifacts. I therefore recommend that the guidelines recognise task contracts, permission scopes, sandbox records, witness events, and human approval gates as useful evidence mechanisms for AI-assisted workflows.

My submission does not claim that my public corpus or any protocol is formally compliant. It proposes a reusable implementation grammar for Article 50: no transparency claim without a system boundary, role classification, trigger identification, disclosure or marking control, evidence artifact, responsible actor, and legal review path.

6. Answer on Article 50(1): direct AI interaction disclosure

Draft answer:

For direct interaction with an AI system, the guidelines should emphasise that the notice must be visible, timely, and understandable to the natural person. It should not be hidden in a general terms-of-service document or depend on the user asking whether the system is AI.

I recommend that implementers record evidence that the notice was actually presented in the relevant interaction flow. A screenshot, versioned UI copy, or interaction-state log can support later review. However, such evidence should not replace the notice itself.

The guidelines should also encourage role clarity. The user should not only know that “AI is involved”, but should understand, at least at a practical level, whether the system is an AI assistant, tool, oracle-like service, automated agent, or AI-assisted workflow with human review.

Shorter version:

Article 50(1) guidance should distinguish the visible user notice from internal evidence that the notice was displayed. A witness or audit log can support review, but it is not a substitute for informing the person in the interaction flow.

7. Answer on Article 50(2): machine-readable marking

Draft answer:

For machine-readable marking, I recommend that the guidelines encourage layered provenance rather than a single fragile mechanism. In practice, embedded metadata may be stripped, watermarks may be altered, and downstream platforms may transform files. Therefore, marking should be supported where possible by sidecar metadata, release manifests, hash references, or other provenance records.

The guidelines should distinguish between three audiences: natural persons, machines/detection systems, and auditors. A visible disclosure informs the person. A machine-readable mark supports detection. A provenance or witness record supports audit. These layers should be connected, but not confused.

For AI-assisted publication workflows, the mark or metadata should ideally identify the artifact, the AI assistance scope, the responsible human/editorial actor, and the review state.

Shorter version:

Machine-readable marking should be treated as part of a provenance chain, not as the whole transparency control. Visible disclosure, metadata, witness/audit evidence, and human responsibility serve different audiences and should be connected.

8. Answer on Article 50(3): emotion recognition and biometric categorisation

Draft answer:

The guidelines should encourage deployers to provide clear positive notices when emotion recognition or biometric categorisation is used, and clear negative declarations where systems may otherwise be misunderstood.

This is important because many systems perform state classification, style adaptation, risk scoring, sentiment analysis, or user-context adaptation without necessarily being designed as legally relevant emotion recognition or biometric categorisation systems. The boundary should be technically described and legally reviewed.

I recommend that the guidelines encourage implementers to maintain a technical function map: what inputs are processed, whether biometric data is used, whether emotional or biometric categories are inferred, who sees the result, and whether the result affects the person.

Shorter version:

The guidelines should help distinguish ordinary state/workflow classification from legally relevant emotion recognition or biometric categorisation, while requiring clear notices where such functions are actually used.

9. Answer on Article 50(4): deepfakes and public-interest publications

Draft answer:

For deepfakes and AI-generated public-interest publications, I recommend stronger guidance on evidence of human editorial responsibility. A general statement that a human reviewed the output is not enough unless the review is substantive, scoped, and attributable.

A useful human review record should identify the reviewed artifact, reviewer or responsible editor, review date, review scope, publication decision, and correction route. This is especially important where human review affects whether a disclosure is required or how it is presented.

For generated or manipulated media, visible disclosure should be combined with machine-readable provenance where possible, because metadata alone may not be visible to the public and may be removed downstream.

The guidelines should also address AI-assisted workflows where CLI/cloud agents prepare publication artifacts, not only cases where a model directly emits final text or media.

Shorter version:

For Article 50(4), human editorial responsibility should be evidenced, not merely declared. Generated media and public-interest AI-assisted publications should combine visible disclosure, provenance/marking, and a recorded human review or correction path.

10. Suggested implementation pattern for the guidelines

Draft answer:

I suggest including an implementation pattern such as:

1. Define the AI system or AI-assisted workflow.
2. Classify the actor role: provider, deployer, publisher/editorial actor, enterprise operator, agent operator, or mixed role.
3. Identify the Article 50 trigger.
4. Define the visible user-facing disclosure.
5. Define the machine-readable mark or provenance mechanism where applicable.
6. Record audit or witness evidence that the control was applied.
7. Record human review and editorial responsibility where relevant.
8. Define a correction, takedown, dispute, or review route.
9. Preserve evidence with data minimisation.
10. Seek legal review for uncertain or high-impact use cases.

This pattern would help implementers move from declarative transparency to testable transparency.

11. Suggested language on evidence

Draft answer:

The guidelines could clarify that evidence of transparency is not the same as transparency itself. For example, a log showing that a notice was displayed can help an auditor, but the person must still have received a clear notice. A machine-readable mark can help detection, but visible disclosure may still be required. A human review record can support editorial responsibility, but only if the review is substantive and attributable.

12. Suggested language on AI-assisted workflows and agents

Draft answer:

I recommend that the guidelines explicitly address AI-assisted workflows involving tool-using or CLI/cloud agents. These systems may create or modify files, generate release artifacts, prepare publications, run validators, or produce metadata. Transparency guidance should therefore consider task contracts, permission scopes, sandbox records, witness events, and human approval gates as useful evidence mechanisms.

This would be particularly useful for SMEs, independent researchers, publishers, developers, and organisations that use AI agents internally before publishing content or systems externally.

13. Suggested language on public-interest AI-generated text

Draft answer:

For AI-generated or AI-assisted text on matters of public interest, the guidelines should clarify what level of human involvement counts as meaningful editorial responsibility. In my view, a superficial spelling pass should not be treated the same as substantive review. The evidence should show who reviewed the content, what was reviewed, what was changed or accepted, and who remains responsible after publication.

14. Suggested language on provenance

Draft answer:

The guidelines should encourage provenance practices that remain useful across formats. For example, Markdown, PDF, HTML, images, audio, video, and code releases may require different marking methods. A practical approach may combine visible labels, embedded metadata, C2PA or equivalent records where feasible, JSON sidecar metadata, release manifests, and hash references.

The key is not to rely on one fragile mechanism, but to preserve a reviewable chain.

15. Suggested language on data minimisation

Draft answer:

The guidelines should also warn against overcollection in the name of transparency. Audit and witness records should preserve the boundary event, provenance, and review state, but should not become a surveillance archive or store unnecessary private prompts, conversations, secrets, or personal data.

A useful principle is: enough evidence for review, minimal raw leakage.

16. Suggested language on mixed roles

Draft answer:

The guidelines could include examples for mixed-role actors. Independent creators, SMEs, researchers, publishers, and enterprise operators may use third-party AI systems, local models, AI agents, and human editors in one workflow. In such cases, a single label such as “provider” or “deployer” may not explain the entire chain. The guidance should encourage role classification per workflow stage.

17. Suggested attachment note

Use if the EUSurvey form allows attachments or supporting links:

I attach or link a short supporting package titled “Article50_Transparency_Submission_v0_1”. It contains a cover note, compliance mapping pack, system-role-obligation-control-evidence table, L4 Witness/provenance bridge, and non-claims/legal review boundary. The package is a technical implementation contribution and does not claim formal compliance or certification.

If only a link is allowed:

Supporting package link: [insert GitHub or Zenodo URL after publication/freeze]

Do not insert a private local path.

18. Suggested closing statement

Draft answer:

The final guideline should help implementers avoid both under-disclosure and overclaiming. Transparency should be concrete enough for users, machine-readable enough for detection, evidenced enough for audit, and responsible enough for legal review. My recommendation is simple: no transparency claim without a system boundary, actor role, trigger, disclosure or marking control, evidence artifact, and responsible human or organisation.

19. Ultra-short final answer

Use where the form gives very limited space:

Please treat Article 50 transparency as an evidence chain: system boundary, actor role, trigger, user notice, machine-readable mark/provenance, audit/witness evidence, human review, responsible actor, and correction route. This prevents decorative labels and supports practical enforcement.

20. A/B slots for tone control

Slot A — formal / official

This contribution proposes a technical evidence-chain model for implementing Article 50 transparency obligations. It does not claim certification or formal compliance. It recommends that every transparency claim be mapped to a concrete system boundary, actor role, obligation trigger, disclosure or marking control, evidence artifact, responsible actor, and legal review path.

Slot B — plain explanation

The practical problem is not whether someone says “AI was used”. The practical problem is whether we can see what system was used, who was responsible, what notice or marking was applied, what evidence exists, and how errors can be corrected.

Use Slot A in the official consultation.

21. Claims to avoid in EUSurvey

Do not write:

My system is Article 50 compliant.
My architecture proves compliance.
L4 Witness solves transparency.
CLI Agent Mesh is a compliance system.
The EU should recognise c as a legal entity.
My corpus already covers all Article 50 obligations.

Write instead:

My submission proposes an architecture-aligned evidence-chain mapping that may help providers and deployers implement Article 50 transparency obligations in a more auditable way.

22. Final checklist before submission

- ☐ Check actual questionnaire fields.
- ☐ Respect character limits.
- ☐ Use official consultation channel.
- ☐ Remove excessive theory.
- ☐ Keep response focused on Article 50.
- ☐ Do not overclaim compliance.
- ☐ Keep CDDP secondary or omit if space is limited.
- ☐ Include CLI Agent Mesh only as workflow provenance/evidence bridge.
- ☐ Include supporting package link only after final public freeze.
- ☐ Human author reviews all final text.
- ☐ Legal review if any concrete compliance claim is added.

23. Recommended final response package

If the form permits one main response and one attachment/link, use:

Main response:

This contribution proposes an evidence-chain approach to Article 50 transparency implementation. Each transparency claim should be mapped from a concrete system boundary and actor role to the relevant obligation trigger, visible disclosure, machine-readable marking or provenance mechanism, audit/witness evidence, human review, responsible actor, and correction route. This is particularly important for AI-assisted publication workflows and tool-using AI agents, where the output may pass through multiple stages before publication. The contribution is architecture-aligned and technical; it does not claim formal compliance or certification for any deployed system.

Attachment/link note:

A short supporting package is available with a compliance mapping table, system-role-obligation-control-evidence matrix, L4 Witness/provenance bridge, and non-claim/legal review boundary.

24. References to cite or link if allowed

Official sources:

1. European Commission, consultation on draft guidelines on transparency obligations under the AI Act:
<https://digital-strategy.ec.europa.eu/en/consultations/consultation-draft-guidelines-transparency-obligations-under-ai-act>
2. European Commission, draft guidelines on implementation of transparency obligations for certain AI systems under Article 50 of the AI Act:
<https://digital-strategy.ec.europa.eu/en/library/draft-guidelines-implementation-transparency-obligations-certain-ai-systems-under-article-50-ai-act>
3. European Commission / AI Office, Code of Practice on marking and labelling of AI-generated content:
<https://digital-strategy.ec.europa.eu/en/policies/code-practice-ai-generated-content>

Supporting technical source:

1. Kotov Ivan, C-Governed CLI Agent Mesh v0.1.1:
<https://github.com/Kot141078/c-governed-cli-agent-mesh>
2. DOI / Zenodo archive for C-Governed CLI Agent Mesh v0.1.1:
<https://doi.org/10.5281/zenodo.20257232>

25. Final control statement

The EUSurvey response should be short enough to be read and strict enough not to overclaim.

The message is:

Make Article 50 transparency auditable:
system, role, trigger, disclosure, marking/provenance, evidence, human responsibility, correction route.

Do not make the message:

Please validate my whole architecture.

That is the boundary.